

Assignment 2 - CS703 Optimisation and Computing

HOANG Thi Linh

April 14, 2025

Question 1

(a) Implement dual solver for LP MDP

```
105 def solve_MDP_LP_dual(states, actions, trans_probs, reward, b_0, gamma):
106     # TODO:create Variables
107     V = cp.Variable()
108     x = cp.Variable((len(states), len(actions)), nonneg=True) # Variables x(s,a)
109     r_sa = np.zeros((len(states), len(actions)))
110     for s in states:
111         for a in actions:
112             r_sa[s, a] = trans_probs[s, a, :] @ reward[s, a, :]
113
114     # objective function = maximize x(s,a)*reward(s,a)
115     objective = cp.Maximize(cp.sum(cp.multiply(x, r_sa)))
116     constraints = []
117     # Constraints
118     # flow constraint sum(x(s',a')) = b_0(s') + gamma * sum(x(s,a))
119     for sprime in states:
120         lhs = cp.sum(x[sprime, :]) # sum_{a'} x(s',a')
121         # sum_{s,a} P(s'|s,a)* x(s,a):
122         flow_in = cp.sum(cp.multiply(trans_probs[:, :, sprime], x))
123         rhs = b_0[sprime] + gamma * flow_in
124         constraints.append(lhs == rhs)
125
126     # TODO:solve the dual LP problem
127     prob = cp.Problem(objective, constraints)
128     prob.solve(solver=cp.GLPK, verbose=False)
129
130     # TODO: extract deterministic policy pi(s)
131     x_opt = x.value
132     pi = [None] * len(states)
133     pi = [np.argmax(x_opt[s, :]) for s in states]
134
135     return prob.value, V.value, pi
```

Figure 1: Screenshot of implementation of solve_MDP_LP_dual

(b)

The optimal objective value of the dual LP is **0.0789681742618707**, which is equal to the primal LP value **0.078968174261873864** up to numerical precision. The small difference (dual gap $\approx 3.16 \times 10^{-16}$) confirms that strong duality holds, as expected for MDPs.

```

*****results from solving primal LP*****

objective value of primal LP: 0.07896817426187386

optimal policy from primal LP:

E      E      E      Goal
N      N      N      Tiger
N      E      N      W

*****results from solving dual LP*****

objective value of dual LP: 0.0789681742618707

optimal policy from dual LP:

E      E      E      Goal
N      E      N      Tiger
N      E      N      W

Dual gap: 3.164135620181696e-15

```

Figure 2: Screenshot of running of solve_MDP_LP_dual

(c) The derivation of dual of the MDP Dual LP

We consider the MDL Dual LP in this question:

$$\begin{aligned}
& \underset{x}{\text{maximize}} && \sum_{s,a} x(s,a) R(s,a) \\
& \text{subject to} && \sum_{a'} x(s',a') = b_0(s') + \gamma \sum_{s,a} P(s'|s,a) x(s,a) \quad \forall s' \\
& && x(s,a) \geq 0 \quad \forall s,a
\end{aligned} \tag{1}$$

Introduce a dual variable $V(s) \in \mathbb{R}$ for each equality constraint at state s , the Lagrangian function associated with (1) is:

$$\begin{aligned}
L(x, V) &= \sum_{s,a} x(s,a) R(s,a) + \sum_{s'} V(s') \left[- \sum_a x(s',a) + b_0(s') + \gamma \sum_{s,a} P(s'|s,a) x(s,a) \right] \\
&= \sum_{s'} V(s') b_0(s') + \sum_{s,a} x(s,a) \left[R(s,a) - V(s) + \gamma \sum_{s'} P(s'|s,a) V(s') \right]
\end{aligned}$$

To ensure $\sup_{x \geq 0} L < +\infty$ each coefficient of $x(s,a)$ must be ≤ 0 :

$$R(s,a) - V(s) + \gamma \sum_{s'} P(s'|s,a) V(s') \leq 0 \Leftrightarrow V(s) \geq R(s,a) + \gamma \sum_{s'} P(s'|s,a) V(s') \quad \forall s,a$$

then the dual of problem (1) is the same as the MDP primal LP as:

$$\begin{aligned}
& \underset{V}{\text{minimize}} && \sum_{s'} V(s') b_0(s') \\
& \text{subject to} && V(s) \geq R(s,a) + \gamma \sum_{s'} P(s'|s,a) V(s') \quad \forall s,a
\end{aligned} \tag{2}$$

Question 2

(a) Formulating optimization problem.

We consider a profit maximization problem (what is the best selling price P to maximize the profit) for a newly designed fashion bag. The setup involves the following parameters:

- A fixed cost of \$700,000 is required for manufacturing setup and marketing.
- Each bag has a production cost of \$110.
- Market demand follows the linear relation: Customer Demand = 70000 - P , with P is the selling price per unit.
- Due to capacity constraints, production is limited to a maximum of 30,000 bags.

Let n denote the number of bags produced and sold. Then, maximizing the total profit as this problem:

$$\begin{aligned} & \underset{n, P}{\text{maximize}} && nP - 110n - 700000 \\ & \text{subject to} && n \leq 70000 - P \\ & && n \leq 30000 \\ & && n \geq 0 \\ & && P \geq 0 \end{aligned} \tag{3}$$

which is equivalent to:

$$\begin{aligned} & \underset{n, P}{\text{minimize}} && f(n, P) = -nP + 110n + 700000 \\ & \text{subject to} && n + P - 70000 \leq 0 \\ & && n - 30000 \leq 0 \\ & && -n \leq 0 \\ & && -P \leq 0 \end{aligned} \tag{4}$$

(b) Solving problem

We will then solve the problem from (4) by introducing the Lagrangian multipliers $\lambda_1, \lambda_2, \lambda_3, \lambda_4$ for constraints and each $\lambda_i \geq 0$. The Lagrangian function associated of this problem is:

$$L(n, P, \lambda) = -nP + 110n + 700000 + \lambda_1(n + P - 70000) + \lambda_2(n - 30000) - \lambda_3n - \lambda_4P \tag{5}$$

The KKT conditions:

- Stationarity:

$$\frac{\partial L}{\partial n} = -P + 110 + \lambda_1 + \lambda_2 - \lambda_3 = 0, \quad \frac{\partial L}{\partial P} = -n + \lambda_1 - \lambda_4 = 0.$$

- Complementary Slackness:

$$\begin{aligned} \lambda_1(n + P - 70000) &= 0, \\ \lambda_2(n - 30000) &= 0, \\ \lambda_3n &= 0, \\ \lambda_4P &= 0. \end{aligned}$$

- Primal feasibility:

$$n + P \leq 70000, n \leq 30000, n \geq 0, P \geq 0.$$

- Feasibility:

$$\lambda_1 \geq 0, \lambda_2 \geq 0, \lambda_3 \geq 0, \lambda_4 \geq 0.$$

Solving the system of KKT conditions we have $n^* = 30000, P^* = 40000, \lambda_1 = 30000, \lambda_2 = 9890, \lambda_3 = \lambda_4 = 0$.

Question 3

(a) Problem Formulation

Variables. Let

$$x_{ij}^k = \begin{cases} 1, & \text{if agent } k \in \{1, 2\} \text{ traverses edge } (i, j) \in E, \\ 0, & \text{otherwise.} \end{cases}$$

Objective. Minimize the total traversal cost: $\min_x \sum_{k=1}^2 \sum_{(i,j) \in E} c_{ij} x_{ij}^k$.

Constraints.

- Source departure: $\sum_{j:(s_k,j) \in E} x_{s_k j}^k = 1, \quad k = 1, 2.$
- Sink arrival: $\sum_{i:(i,d_k) \in E} x_{i d_k}^k = 1, \quad k = 1, 2.$
- Flow conservation: $\sum_{i:(i,v) \in E} x_{i v}^k - \sum_{j:(v,j) \in E} x_{v j}^k = 0, \quad k = 1, 2, \quad v \neq s_k, d_k.$
- Coverage: $\sum_{k=1}^2 \sum_{i:(i,v) \in E} x_{i v}^k \geq 1, \quad \forall v \in V.$
- Binary: $x_{ij}^k \in \{0, 1\}, \quad k = 1, 2, (i, j) \in E.$

The optimization problem is

$$\begin{aligned} \min_x \quad & \sum_{k=1}^2 \sum_{(i,j) \in E} c_{ij} x_{ij}^k \\ \text{s.t.} \quad & \sum_{j:(s_k,j) \in E} x_{s_k j}^k = 1, & k = 1, 2, \\ & \sum_{i:(i,d_k) \in E} x_{i d_k}^k = 1, & k = 1, 2, \\ & \sum_{i:(i,v) \in E} x_{i v}^k - \sum_{j:(v,j) \in E} x_{v j}^k = 0, & k = 1, 2, \quad v \neq s_k, d_k, \\ & \sum_{k=1}^2 \sum_{i:(i,v) \in E} x_{i v}^k \geq 1, & \forall v \in V, \\ & x_{ij}^k \in \{0, 1\}, & k = 1, 2, (i, j) \in E. \end{aligned}$$

Computational tractability. This is a mixed-integer program with binary routing and global coverage constraints and is NP-hard, so it is not solvable in polynomial time in the worst case. Therefore, the problem is not tractable for large graphs.

(b) Dual Decomposition.

Relax the coverage constraints $\sum_k \sum_{i:(i,v) \in E} x_{iv}^k \geq 1$ with multipliers $\lambda_v \geq 0$. The Lagrangian separates by agent:

$$\begin{aligned} L(x, \lambda) &= \sum_{k=1}^2 \sum_{(i,j) \in E} c_{ij} x_{ij}^k + \sum_{v \in V} \lambda_v \left(1 - \sum_{k=1}^2 \sum_{i:(i,v) \in E} x_{iv}^k \right) \\ &= \sum_{v \in V} \lambda_v + \sum_{k=1}^2 \underbrace{\sum_{(i,j) \in E} (c_{ij} - \lambda_j) x_{ij}^k}_{\Phi_k(x^k; \lambda)}. \end{aligned}$$

Hence the dual is

$$\max_{\lambda \geq 0} \left[\sum_{v \in V} \lambda_v + \sum_{k=1}^2 \min_{x^k \in \mathcal{X}_k} \Phi_k(x^k; \lambda) \right],$$

where each subproblem for agent k is

$$\min_{x^k \in \mathcal{X}_k} \sum_{(i,j) \in E} (c_{ij} - \lambda_j) x_{ij}^k,$$

and \mathcal{X}_k encodes that agent's flow-conservation and integrality.

(c) Tractability of Subproblems.

Each subproblem is a *shortest-path* or *min-cost-flow* problem with nonnegative modified costs $c_{ij} - \lambda_j$. Such problems can be solved in polynomial time, so each subproblem is tractable.

Question 4

Consider the problem

$$\begin{aligned} \min_{x_1, x_2, x_3 > 0} \quad & \max \left\{ \frac{x_1}{x_2}, \frac{\sqrt{x_3}}{x_2} \right\} \\ \text{subject to} \quad & x_1^2 + \frac{2x_2}{x_3} \leq \sqrt{x_2} \\ & \frac{x_1}{x_2} \geq x_3^2 \end{aligned} \tag{6}$$

(a) Show that the above problem can convert into a convex optimization problem.

Following [1], the standard form of a geometric programming is:

$$\begin{aligned} \text{minimize} \quad & f_0(x) \\ \text{subject to} \quad & f_i(x) \leq 1, \quad i = 1, \dots, m \\ & g_i(x) = 1, \quad i = 1, \dots, p \end{aligned}$$

where f_i is posynomial functions, g_i are mononials, and x_i are the optimization variables. There is an implicit constraint that the variables are positive.

Introduce $t > 0$ such that $\frac{x_1}{x_2} \leq t$, $\frac{\sqrt{x_3}}{x_2} \leq t$, we can convert the problem (6) to a geometric programming:

$$\begin{aligned} \text{minimize} \quad & t \\ \text{subject to} \quad & x_1 x_2^{-1} t^{-1} \leq 1, \\ & x_3^{1/2} x_2^{-1} t^{-1} \leq 1, \\ & x_1^2 x_2^{-1/2} + 2 x_2^{1/2} x_3^{-1} \leq 1, \\ & x_3^2 x_2 x_1^{-1} \leq 1. \end{aligned} \tag{7}$$

Set $y_i = \log x_i$ ($i = 1, 2, 3$) and $u = \log t$, the problem (7) is equivalent to the convex optimization:

$$\begin{aligned}
 & \underset{y, u}{\text{minimize}} && u \\
 & \text{subject to} && y_1 - y_2 - u \leq 0, \\
 & && \frac{1}{2}y_3 - y_2 - u \leq 0, \\
 & && \log\left(\exp(2y_1 - \frac{1}{2}y_2) + 2 \exp(\frac{1}{2}y_2 - y_3)\right) \leq 0, \\
 & && 2y_3 + y_2 - y_1 \leq 0.
 \end{aligned} \tag{8}$$

(b) Solve the resulting convex problem using cvxpy

The implementation of GP problem as the screenshot in Fig 3. Solving the problem we get the solution in the (y, u) -space and the corresponding original (x, t) values are:

$$\begin{aligned}
 y_1^* &= -1.3785, & x_1^* &= e^{y_1^*} \approx 0.2520, \\
 y_2^* &= -1.5290, & x_2^* &= e^{y_2^*} \approx 0.2168, \\
 y_3^* &= 0.0753, & x_3^* &= e^{y_3^*} \approx 1.0782, \\
 u^* &= 1.5666, & t^* &= e^{u^*} \approx 4.7900.
 \end{aligned}$$

```

question4.py
1  import cvxpy as cp
2  import numpy as np
3
4  # Variables
5  y1 = cp.Variable(name="y1")
6  y2 = cp.Variable(name="y2")
7  y3 = cp.Variable(name="y3")
8  u = cp.Variable(name="u")
9
10 # Constraints
11 constraints = [
12     y1 - y2 - u <= 0, # y1 - y2 - u <= 0
13     0.5*y3 - y2 - u <= 0, # (1/2) y3 - y2 - u <= 0
14     # log(exp(2y1-0.5y2) + 2 exp(0.5y2 - y3)) <= 0
15     cp.log_sum_exp(cp.hstack([
16         2*y1 - 0.5*y2,
17         cp.log(2) + 0.5*y2 - y3
18     ])) <= 0,
19     2*y3 + y2 - y1 <= 0 # 2y3 + y2 - y1 <= 0
20 ]
21
22 obj = cp.Minimize(u) # objective
23
24 # Solve
25 prob = cp.Problem(obj, constraints)
26 prob.solve()
27
28 # Convert y back to x via x_i = exp(y_i)
29 x1_val, x2_val, x3_val = np.exp(y1.value), np.exp(y2.value), np.exp(y3.value)
30
31 print(f"Optimal y1 = {y1.value:.4f}, y2 = {y2.value:.4f}, y3 = {y3.value:.4f}, u = {u.value:.4f}")
32 print(f"Corresponding x1 = exp(y1) = {x1_val:.4f}")
33 print(f"Corresponding x2 = exp(y2) = {x2_val:.4f}")
34 print(f"Corresponding x3 = exp(y3) = {x3_val:.4f}")
35

```

Figure 3: Screenshot of the implementation using cvxpy for the problem (8)

Question 5

On the 4×4 grid we single out three neighbouring nodes i – j – ℓ with i immediately left of j and ℓ immediately above j .

(a) Random-variable domains

For every location $v \in \{i, j, \ell\}$ define a variable

$$X_v \in \{T, D, L, R\},$$

where T = top, D = down, L = left, R = right. Thus X_v records the *single* edge that agent v chooses to scan.

(b) Pairwise potential functions

A reward is obtained only when *both* endpoints of a target edge scan *toward* each other; otherwise the reward is 0.

Edge (j, i) (horizontal target). Reward $t_{ji} > 0$ is earned only when j scans *left* (L) and i scans *right* (R). Hence

$$\theta_{j,i}(X_j, X_i) = \begin{cases} t_{ji}, & (X_j, X_i) = (L, R), \\ 0, & \text{otherwise.} \end{cases}$$

and potential entries in table form is:

X_j	X_i	$\theta_{j,i}(X_j, X_i)$
T	T	0
T	D	0
T	L	0
T	R	0
D	T	0
D	D	0
D	L	0
D	R	0
L	T	0
L	D	0
L	L	0
L	R	t_{ji}
R	T	0
R	D	0
R	L	0
R	R	0

Similarly, **Edge (j, ℓ) (vertical target).** Reward $t_{j\ell} > 0$ is earned only when j scans *top* (T) and ℓ scans *down* (D):

$$\theta_{j,\ell}(X_j, X_\ell) = \begin{cases} t_{j\ell}, & (X_j, X_\ell) = (T, D), \\ 0, & \text{otherwise.} \end{cases}$$

with table form is:

X_j	X_ℓ	$\theta_{j,\ell}(X_j, X_\ell)$
T	T	0
T	D	$t_{j\ell}$
T	L	0
T	R	0
D	T	0
D	D	0
D	L	0
D	R	0
L	T	0
L	D	0
L	L	0
L	R	0
R	T	0
R	D	0
R	L	0
R	R	0

These potentials ensure that a reward is collected *iff* both agents involved in a target edge coordinate their scan directions toward that edge; any other combination yields zero.

References

- [1] Stephen Boyd et al. “A tutorial on geometric programming”. In: *Optimization and engineering* 8 (2007), pp. 67–127.